

TITANIC: SURVIVAL DYNAMICS & PREDICTIVE MODELING

The objective of this notebook is to go beyond simple descriptive statistics and uncover the hidden dynamics of the Titanic disaster. We aim to quantify the determining factors of survival to answer one central question: *Which passenger profiles were spared from the tragedy, and what underlying variables tipped the scales in their favor?*

To unravel this, we will follow a multi-step analytical framework:

- **Environment Setup:** Preparing our tools and libraries.
- **Data Loading & Cleaning:** Consolidating the dataset and addressing missing values.
- **Univariate Profiling:** Understanding the individual characteristics of the passengers.
- **Bivariate Survival Analysis:** Examining how specific features directly impacted survival rates.
- **Multivariate Analysis:** Uncovering complex interactions between multiple variables.
- **Predictive Modeling:** Building robust machine learning algorithms to forecast survival probabilities.

1. Environment setup

```
In [88]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import re
import plotly.express as px
import plotly.graph_objects as go
from plotly.subplots import make_subplots
from scipy import stats
from scipy.stats import chi2_contingency, ttest_ind, f_oneway, mannwhitneyu
from sklearn.model_selection import StratifiedKFold, cross_val_score
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier, VotingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
import warnings
import warnings, json, math, textwrap
warnings.filterwarnings("ignore")

sns.set_theme(style="whitegrid", context="talk")
plt.rcParams["figure.figsize"] = (12, 6)
plt.rcParams["axes.titlesize"] = 14
plt.rcParams["axes.titleweight"] = "bold"
PALETTE = {"Survived": "#2ecc71", "Died": "#e74c3c", "Neutral": "#3498db"}
```

```
pd.set_option("display.max_columns", 50)
pd.set_option("display.width", 160)
pd.set_option("display.float_format", lambda x: f"{x:,.3f}")

RANDOM_STATE = 42
print("Environment ready")
```

Environment ready

2. Data loading & Missing-value strategy

2.1. Data loading

We load both `train.csv` and `test.csv`.

```
In [89]: train = pd.read_csv("/kaggle/input/titanic/train.csv")
        test  = pd.read_csv("/kaggle/input/titanic/test.csv")

        # Tag origin so we can later concatenate safely
        train["__set"] = "train"
        test["__set"]  = "test"
        test["Survived"] = np.nan

        full = pd.concat([train, test], ignore_index=True, sort=False)
        print(f"Train shape: {train.shape} | Test shape: {test.shape} | Combined: {full.sha
```

Train shape: (891, 13) | Test shape: (418, 13) | Combined: (1309, 13)

```
In [90]: train.head(5)
```

Out[90]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.000	1	0	A/5 21171	7.25
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.000	1	0	PC 17599	71.28
2	3	1	3	Heikkinen, Miss. Laina	female	26.000	0	0	STON/O2. 3101282	7.92
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.000	1	0	113803	53.10
4	5	0	3	Allen, Mr. William Henry	male	35.000	0	0	373450	8.05



In [91]:

```
test.head(5)
```

Out[91]:


	PassengerId	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	892	3	Kelly, Mr. James	male	34.500	0	0	330911	7.829	NaN	
1	893	3	Wilkes, Mrs. James (Ellen Needs)	female	47.000	1	0	363272	7.000	NaN	
2	894	2	Myles, Mr. Thomas Francis	male	62.000	0	0	240276	9.688	NaN	
3	895	3	Wirz, Mr. Albert	male	27.000	0	0	315154	8.662	NaN	
4	896	3	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	female	22.000	1	1	3101298	12.287	NaN	



2.2. Missing-value strategy

```
In [92]: # Missing-value heatmap
fig, ax = plt.subplots(figsize=(12,5))
sns.heatmap(full.isna(), cbar=False, yticklabels=False, cmap="magma", ax=ax)
ax.set_title("Missing-Value Map (yellow = missing)")
plt.tight_layout(); plt.show()

# Tabular summary
(full.isna().sum().to_frame("missing")
 .assign(pct=lambda d: (d["missing"]/len(full)*100).round(2))
 .query("missing > 0").sort_values("missing", ascending=False))
```



```
Out[92]:
```

	missing	pct
Cabin	1014	77.460
Survived	418	31.930
Age	263	20.090
Embarked	2	0.150
Fare	1	0.080

```
In [93]: # Engineering a Title column
full["Title"] = full["Name"].str.extract(r",\s*([^\.]+)\.")
full["Title"] = full["Title"].str.strip()

# Consolidate rare titles
rare = full["Title"].value_counts()[full["Title"].value_counts() < 10].index
full["Title"] = full["Title"].replace(rare, "Rare")
full["Title"] = full["Title"].replace({"Mlle": "Miss", "Ms": "Miss", "Mme": "Mrs"})
full["Title"].value_counts()
```

```
Out[93]: Title
Mr      757
Miss    260
Mrs     197
Master   61
Rare     34
Name: count, dtype: int64
```

```
In [94]: # Impute Age with group median
full["Age"] = full.groupby(["Pclass", "Sex", "Title"])["Age"].transform(lambda s: s.fillna(s.median()))

# Fallback if any NaN remains
full["Age"] = full["Age"].fillna(full["Age"].median())

# Impute Fare & Embarked
full["Fare"] = full.groupby("Pclass")["Fare"].transform(lambda s: s.fillna(s.median()))
full["Embarked"] = full["Embarked"].fillna(full["Embarked"].mode()[0])

# Cabin → Deck
full["Deck"] = full["Cabin"].astype(str).str[0].replace("n", "U") # U = Unknown
full["HasCabin"] = full["Cabin"].notna().astype(int)

# Family features
full["FamilySize"] = full["SibSp"] + full["Parch"] + 1
full["IsAlone"] = (full["FamilySize"] == 1).astype(int)

# Tickets / counts
full["TicketPrefix"] = (full["Ticket"].str.replace(r"^[A-Za-z]", "", regex=True)
                        .replace("", "NONE"))

print("Imputation & feature engineering complete")
print("Remaining missing values per column:")
print(full.isna().sum()[full.isna().sum() > 0])
```

```
Imputation & feature engineering complete
Remaining missing values per column:
Survived      418
Cabin         1014
dtype: int64
```

3. Univariate profiling

Re-split engineered data back to train for supervised analysis.

```
In [95]: df = full[full["__set"]=="train"].copy()
df["Survived_lbl"] = df["Survived"].map({0: "Died", 1: "Survived"})
print(f"Working analysis frame: {df.shape}")
```

```
Working analysis frame: (891, 20)
```

Note: Although the Titanic historically carried approximately 1,317 passengers and 885 crew members, this exploratory analysis focuses on the 891 passengers in the training dataset (`train.csv`). The other 418 passengers (`test.csv`) are excluded because their survival status is unknown (reserved for predictive modeling). The crew is


not included in this database. The conclusions drawn here are therefore representative of a highly reliable sample (68% of the actual passenger manifest).

3.1. Target distribution

```
In [96]: fig, ax = plt.subplots(1,2, figsize=(13,5))
counts = df["Survived_lbl"].value_counts()
ax[0].pie(counts, labels=counts.index, autopct="%1.1f%%",
          colors=[PALETTE["Died"],PALETTE["Survived"]], startangle=90, wedgeprops=d
ax[0].set_title("Survival Share")

sns.countplot(data=df, x="Survived_lbl", palette=[PALETTE["Died"],PALETTE["Survived"]])
for i,v in enumerate(counts): ax[1].text(i, v+5, str(v), ha="center", fontweight="b")
ax[1].set_title("Survival Counts"); ax[1].set_ylabel("Passengers")
plt.tight_layout(); plt.show()

overall_rate = df["Survived"].mean()
print(f"Overall survival rate: {overall_rate:.1%} ({df['Survived'].sum()} / {len(df)})")
```



Overall survival rate: 38.4% (342.0 / 891)

Interpretation: Of 891 training passengers, **38.4% survived**.

3.2. Categorical feature distributions

```
In [97]: cat_cols = ["Pclass", "Sex", "Embarked", "Title", "Deck", "FamilySize"]
fig, axes = plt.subplots(2,3, figsize=(18,10))

total = len(df)

for i, ax in enumerate(axes.ravel()):
    if i < len(cat_cols):
        col = cat_cols[i]
        order = df[col].value_counts().index
        sns.countplot(data=df, x=col, order=order, ax=ax, palette="viridis")
        ax.set_title(f"Distribution of {col}")
        ax.tick_params(axis="x", rotation=30)


    for p in ax.patches:
```

```

count = int(p.get_height())
pct = 100 * count / total
ax.annotate(f"{count}\n({pct:.1f}%)",
            (p.get_x() + p.get_width() / 2., p.get_height()),
            ha="center", va="bottom", fontsize=15)

plt.tight_layout(); plt.show()

```



Interpretation:

- **Pclass:** 3rd-class dominates. Titanic was economically bottom-heavy.
- **Sex:** 64.7% male.
- **Embarked:** Southampton (S - 72.44%) is by far the dominant port. Note that C refers to Cherbourg, and Q to Queenstown.
- **Title:** "Mr" is the largest group; "Rare" includes honorifics like Dr, Rev, Col, Capt, Lady.
- **Deck:** Unknown deck (U) is the largest bucket due to cabin sparsity.
- **FamilySize:** Solo travellers dominate, with a long tail of large families.

Note: The Titanic had nine metal decks (A through G) plus the upper decks. Here's what they correspond to, from top (closest to the lifeboats) to bottom (closest to the hull):


- *T: Boat Deck. This is where the lifeboats were located. Very few luxury cabins were found here.*
- *A: Promenade Deck. Reserved for First Class.*
- *B, C, D, E: These are the main decks. Decks B, C, and D housed the most luxurious First Class cabins and suites (in the center of the ship). Deck D ("Saloon Deck") also contained the grand First Class dining room. Deck E ("Upper Deck") marked the transition, with First, Second, and Third Class cabins, as well as crew quarters.*
- *F: Middle Deck. Primarily 2nd and 3rd class, plus the galleys.*

- *G: Lower Deck. The lowest deck accessible to passengers. Almost exclusively 3rd class (located very close to the waterline and the boiler room crews).*
- *And the "U"? In our code, we created the category U (Unknown) for all passengers who did not have a cabin number recorded in the dataset (the `cabin` column was empty 77.1% of the time, as shown in the graphic).*

3.3. Numeric feature distributions

```
In [98]: num_cols = ["Age", "Fare", "SibSp", "Parch"]
fig, axes = plt.subplots(2,2, figsize=(16,10))
for ax, col in zip(axes.ravel(), num_cols):
    sns.histplot(df[col].dropna(), kde=True, color="#3498db", bins=40, ax=ax)
    ax.axvline(df[col].median(), color="orange", linestyle="--", label=f"median={df[col].median():.2f}")
    ax.axvline(df[col].mean(), color="red", linestyle="--", label=f"mean={df[col].mean():.2f}")
    ax.set_title(f"{col} distribution"); ax.legend()
plt.tight_layout(); plt.show()

df[num_cols].describe().T
```



```
Out[98]:
```

	count	mean	std	min	25%	50%	75%	max
Age	891.000	29.205	13.553	0.420	21.000	26.000	36.750	80.000
Fare	891.000	32.204	49.693	0.000	7.910	14.454	31.000	512.329
SibSp	891.000	0.523	1.103	0.000	0.000	0.000	1.000	8.000
Parch	891.000	0.382	0.806	0.000	0.000	0.000	0.000	6.000


Interpretation:

- **Age:** Peak around 20-30, plus a children bump under 10. Mean = 29.20 (after imputation). Right-skew is mild.
- **Fare:** Heavily right-skewed (median 14.45 vs mean 32.20). A few first-class suites exceeded 500.
- **SibSp / Parch:** Mostly 0 (solo) or small (1,2, 3, 4). Tail values up to 8 represent large families. They had historically low survival (For example, Sage family lost all 11).

4. Bivariate survival analysis

4.1. Survival by sex (the headline gap)

```
In [99]: g = df.groupby("Sex")["Survived"].agg(["mean", "count"])
g["mean"].plot(kind="bar", color=["#e74c3c", "#2ecc71"], figsize=(7,5))
plt.title("Survival Rate by Sex"); plt.ylabel("Survival rate"); plt.ylim(0,1)
for i,(m,c) in enumerate(zip(g["mean"], g["count"])):
    plt.text(i, m+0.02, f"{m:.1%}\n(n={c})", ha="center", fontweight="bold")
plt.tight_layout(); plt.show()
g
```



```
Out[99]:
```

	mean	count
Sex		
female	0.742	314
male	0.189	577


Interpretation: Females survived at **74.2%** vs males at **18.9%** (a **55-point gender gap**). The "women and children first" protocol visibly dominated outcomes.

4.2 Survival by passenger class

```
In [100...] fig, ax = plt.subplots(1,2, figsize=(14,5))
class_tab = df.groupby("Pclass")["Survived"].mean()
class_tab.plot(kind="bar", color=["#1abc9c", "#9b59b6", "#e67e22"], ax=ax[0])
ax[0].set_title("Survival Rate by Pclass"); ax[0].set_ylabel("Rate"); ax[0].set_yli
for i,v in enumerate(class_tab): ax[0].text(i, v+0.02, f"{v:.1%}", ha="center", fon

sns.countplot(data=df, x="Pclass", hue="Survived_lbl",
              palette=[PALETTE["Died"], PALETTE["Survived"]], ax=ax[1])
ax[1].set_title("Class Composition vs Outcome")
plt.tight_layout(); plt.show()

df.groupby("Pclass")["Survived"].agg(["mean", "count", "sum"])
```




```
Out[100...]
      mean  count  sum
Pclass
1    0.630    216 136.000
2    0.473    184  87.000
3    0.242    491 119.000
```

Interpretation: Survival declines monotonically with class: **1st = 63%, 2nd = 47%, 3rd = 24%**.

4.3. Sex x Class interaction

```
In [101...] pivot = df.pivot_table(index="Sex", columns="Pclass", values="Survived", aggfunc="m
sns.heatmap(pivot, annot=True, fmt=".1%", cmap="RdYlGn", vmin=0, vmax=1, cbar_kws={
plt.title("Survival Rate - Sex x Pclass"); plt.tight_layout(); plt.show()
pivot
```



Out[101...

	Pclass		
	1	2	3
Sex			
female	0.968	0.921	0.500
male	0.369	0.157	0.135

Interpretation: The interaction is striking:



- 1st & 2nd-class **females**: 96.8% & 92.1% survival - near-universal rescue.
- 3rd-class **females**: only 50% - class overrode gender for poorer women.
- 1st-class **males** (36.9%) still out-survived 3rd-class males (13.5%).

4.4. Survival by age

In [102...

```
fig, ax = plt.subplots(1,2, figsize=(15,5))
sns.kdeplot(data=df, x="Age", hue="Survived_lbl", fill=True, common_norm=False,
            palette=[PALETTE["Died"], PALETTE["Survived"]], ax=ax[0])
ax[0].set_title("Age Density by Outcome")

age_band = pd.cut(df["Age"], bins=[0,5,12,18,30,45,60,100],
                  labels=["0-5", "6-12", "13-18", "19-30", "31-45", "46-60", "60+"])
band_rate = df.assign(AgeBand=age_band).groupby("AgeBand")["Survived"].mean()
band_rate.plot(kind="bar", color="#3498db", ax=ax[1])
ax[1].set_title("Survival Rate by Age Band"); ax[1].set_ylim(0,1); ax[1].set_ylabel("Survival Rate")
for i,v in enumerate(band_rate): ax[1].text(i, v+0.02, f"{v:.1%}", ha="center", fontweight="bold")
plt.tight_layout(); plt.show()
band_rate
```





```
Out[102...] AgeBand
0-5      0.705
6-12     0.379
13-18    0.476
19-30    0.290
31-45    0.439
46-60    0.402
60+      0.227
Name: Survived, dtype: float64
```

Interpretation: Children 0-5 had the highest survival (70.5%), validating the "children first" protocol. Seniors (60+) had the lowest (22.7%), possibly due to reduced mobility. The 19-30 male bulge drove the largest absolute loss.

4.5. Survival by fare (log-scaled)

```
In [103...] fig, ax = plt.subplots(1,2, figsize=(15,5))
sns.boxplot(data=df, x="Survived_lbl", y="Fare",
            palette=[PALETTE["Died"],PALETTE["Survived"]], ax=ax[0], showfliers=False)
ax[0].set_title("Fare distribution by outcome")

fare_q = pd.qcut(df["Fare"], 5, labels=["Q1","Q2","Q3","Q4","Q5"])
fq_rate = df.assign(FareQ=fare_q).groupby("FareQ")["Survived"].mean()
fq_rate.plot(kind="bar", color="#16a085", ax=ax[1])
ax[1].set_title("Survival by Fare Quintile"); ax[1].set_ylim(0,1)
for i,v in enumerate(fq_rate): ax[1].text(i, v+0.02, f"{v:.1%}", ha="center", fontw
plt.tight_layout(); plt.show()
fq_rate
```



```
Out[103... FareQ
Q1    0.218
Q2    0.201
Q3    0.424
Q4    0.444
Q5    0.642
Name: Survived, dtype: float64
```

Interpretation: This graph compares the distribution of prices paid by those who survived and those who died. It demonstrates a strong monotonic relationship between economic status and the chances of survival. Survivors generally paid significantly more for tickets than the victims. By dividing the prices into quintiles, a clear gradient emerges where the survival rate increases from approximately 21.8% for the cheapest tickets (Q1) to nearly 64.2% for the most luxurious (Q5). Ticket price thus acts as a continuous proxy for social class and cabin location on the ship. Indeed, higher fares bought access to the upper decks, located in the immediate vicinity of the lifeboats, providing a decisive evacuation advantage.

4.6. Survival by embarkation port

```
In [104... port_rate = df.groupby("Embarked")["Survived"].agg(["mean", "count"])
fig, ax = plt.subplots(1, 3, figsize=(22, 6))
port_rate["mean"].loc[["S", "C", "Q"]].plot(kind="bar", color="#8e44ad", ax=ax[0])
ax[0].set_title("Survival by Port")
ax[0].set_ylim(0, 1)
ax[0].set_ylabel("Rate")
ax[0].tick_params(axis='x', rotation=0)
for i, (p, v) in enumerate(port_rate.loc[["S", "C", "Q"]].iterrows()):
    ax[0].text(i, v["mean"] + 0.02, f'{v["mean"]:.1%}\n(n={int(v["count"])})',
               ha="center", fontweight="bold")


port_class = df.groupby(["Embarked", "Pclass"]).size().unstack().loc[["S", "C", "Q"])
port_class_pct = port_class.div(port_class.sum(axis=1), axis=0)
port_class_pct.plot(kind="bar", stacked=True, ax=ax[1], colormap="Set2")
ax[1].set_title("Class Mix by Port")
ax[1].set_ylabel("Share")
ax[1].tick_params(axis='x', rotation=0)
for c in ax[1].containers:
    labels = [f'{v.get_height()*100:.1f}%' if v.get_height() > 0.05 else '' for v in c]
    ax[1].bar_label(c, labels=labels, label_type='center', color='black', fontweight='bold')
ax[1].legend(title="Pclass", bbox_to_anchor=(1.05, 1), loc='upper left')

pivot_survival = df.pivot_table(index="Embarked", columns="Pclass", values="Survived")
pivot_survival.plot(kind="bar", ax=ax[2], colormap="viridis")
ax[2].set_title("Survival Rate by Class & Port")
ax[2].set_ylabel("Survival Rate")
ax[2].set_ylim(0, 1)
ax[2].tick_params(axis='x', rotation=0)

for container in ax[2].containers:
    labels = [f'{v.get_height():.1%}' if not pd.isna(v.get_height()) else '' for v in container]
    ax[2].bar_label(container, labels=labels, label_type='edge', padding=3, fontweight='bold')
ax[2].legend(title="Pclass", bbox_to_anchor=(1.05, 1), loc='upper left')
```

```
plt.tight_layout()
plt.show()

port_rate
```



Out[104...

	mean	count
Embarked		
C	0.554	168
Q	0.390	77
S	0.339	646

Interpretation:



- Cherbourg (C) had the highest survival (55.4%), but this is confounded by class: Cherbourg had a disproportionate share of 1st-class passengers (50.6%).
- Southampton (S) is mostly 3rd-class (93.5%), explaining its 33.9% survival.
- In general, for third class, the survival rate is low everywhere, but it varies slightly depending on the port of embarkation. This confirms that the port is merely a proxy (a surrogate variable): it is the class (and not the port itself) that dictates the chances of survival.

4.7. Family size

In [105...

```
fam_rate = df.groupby("FamilySize")["Survived"].agg(["mean", "count"])
fig, ax = plt.subplots(1, 2, figsize=(14, 5))
fam_rate["mean"].plot(kind="bar", color="#d35400", ax=ax[0])
ax[0].set_title("Survival by Family Size"); ax[0].set_ylim(0, 1); ax[0].set_ylabel("Survival Rate")
for i, (s, v) in enumerate(fam_rate.iterrows()):
    ax[0].text(i, v["mean"] + 0.02, f"{v['mean']:.1%}\n(n={int(v['count'])})", ha="center")

df.groupby("FamilySize")["Survived"].count().plot(kind="bar", color="#f78c8d", ax=ax[1])
ax[1].set_title("Family Size Counts")
plt.tight_layout(); plt.show()
fam_rate
```



Out[105...

	mean	count
FamilySize		
1	0.304	537
2	0.553	161
3	0.578	102
4	0.724	29
5	0.200	15
6	0.136	22
7	0.333	12
8	0.000	6
11	0.000	7


Interpretation: Solo travellers (30.4%) and families ≥ 5 fared worst, while families of 2-4 had 50-70% survival.

5. Multivariate & Segment interaction analysis

5.1. Survival by Class x Family size

In [106...

```
pv = df.pivot_table(index="Pclass", columns="FamilySize", values="Survived", aggfun
sns.heatmap(pv, annot=True, fmt=".1%", cmap="RdYlGn", vmin=0, vmax=1, cbar_kws={"la
plt.title("Survival - Pclass x FamilySize"); plt.tight_layout(); plt.show()
```



Interpretation: This distribution shows some evidence. 1st-class family-of-5 cell (100% survival) shows that small wealthy groups essentially all made it. Conversely, 3rd-class large families form had a very weak survival rate.

5.2. Age x Sex survival surface


```
In [107... g = sns.FacetGrid(df, col="Sex", hue="Survived_lbl", height=5, palette=[PALETTE["Di
g.map_dataframe(sns.kdeplot, x="Age", fill=True, common_norm=False, alpha=0.5)
g.add_legend(); g.set_titles("{col_name}")
g.fig.suptitle("Age x Sex outcome density", y=1.04, fontsize=14, fontweight="bold")
plt.show()
```



Interpretation For males, the survival density peaks in childhood then collapses through adulthood. For females, survival dominates across all ages - confirming that "children first" was applied on top of "women first", not as a substitute.

5.3. Title x Pclass survival grid

```
In [108... pv2 = df.pivot_table(index="Title", columns="Pclass", values="Survived", aggfunc="m
sns.heatmap(pv2, annot=True, fmt=".1%", cmap="RdYlGn", vmin=0, vmax=1)
plt.title("Survival - Title x Pclass"); plt.tight_layout(); plt.show()
```



Interpretation: "Mrs" and "Miss" titles in 1st/2nd class (90-95%). The protocol was near-absolute for high-status women. "Master" (boys) also performed well (79.76% overall).

6. Feature engineering

```
In [109... engineered = df.copy()
engineered["AgeBand"] = pd.cut(engineered["Age"], bins=[0,5,12,18,30,45,60,100],
                                labels=["0-5", "6-12", "13-18", "19-30", "31-45", "46-
engineered["FareBand"] = pd.qcut(engineered["Fare"], 4, labels=["Low", "Mid", "High"]
engineered["FamilyType"] = pd.cut(engineered["FamilySize"], [-1,1,4,11],
                                   labels=["Solo", "Small", "Large"])


# Quick KPI panel
panel = engineered.groupby(["Pclass", "Sex", "FamilyType"])["Survived"].agg(["mean", "
panel.head(20)
```

Out[109...

			mean	count
Pclass	Sex	FamilyType		
1	female	Solo	0.971	34
		Small	0.964	56
		Large	1.000	4
	male	Solo	0.333	75
		Small	0.444	45
		Large	0.000	2
2	female	Solo	0.906	32
		Small	0.929	42
		Large	1.000	2
	male	Solo	0.097	72
		Small	0.278	36
		Large	NaN	0
3	female	Solo	0.617	60
		Small	0.561	57
		Large	0.111	27
	male	Solo	0.121	264
		Small	0.250	56
		Large	0.037	27

In [110...

```
# Deck survival
df["Deck"] = df["Cabin"].fillna("U").astype(str).str[0]
deck_rate = df.groupby("Deck")["Survived"].agg(["mean", "count"]).query("count>=5").
deck_rate.plot(kind="bar", y="mean", legend=False, color="#27ae60", figsize=(10,5))
plt.title("Survival by Deck (n≥5)"); plt.ylim(0,1); plt.ylabel("Survival rate")
for i,v in enumerate(deck_rate["mean"]): plt.text(i, v+0.02, f"{v:.1%}", ha="center")
plt.tight_layout(); plt.show()
deck_rate
```



Out[110...

	mean	count
Deck		
D	0.758	33
E	0.750	32
B	0.745	47
F	0.615	13
C	0.593	59
A	0.467	15
U	0.300	687

Interpretation: Decks B-D (1st class, boat-deck adjacent) show >75% survival. The U (unknown cabin) bucket was 30%, confirming missing cabins correlate with poorer passengers.

7. Statistical Validation

We now formally test whether observed differences are statistically significant.

```
In [111... def chi2_test(col):
    ct = pd.crosstab(df[col], df["Survived"])
    chi2, p, dof, _ = chi2_contingency(ct)
    return chi2, p, dof

cat_features = ["Pclass", "Sex", "Embarked", "Title", "Deck", "FamilySize", "IsAlone", "HasSiblings"]
results = []
for c in cat_features:
    chi2, p, dof = chi2_test(c)
    results.append({"Feature":c, "Chi2":chi2, "dof":dof, "p-value":p,
```

```

        "Significant": "✅" if p<0.05 else "❌"})
chi_df = pd.DataFrame(results).sort_values("Chi2", ascending=False)
chi_df

```

Out[111...

	Feature	Chi2	dof	p-value	Significant
3	Title	283.312	4	0.000	✅
1	Sex	260.717	1	0.000	✅
0	Pclass	102.889	2	0.000	✅
4	Deck	99.164	8	0.000	✅
7	HasCabin	87.941	1	0.000	✅
5	FamilySize	80.672	8	0.000	✅
6	IsAlone	36.001	1	0.000	✅
2	Embarked	25.964	2	0.000	✅

Interpretation: All categorical features are highly significant ($p < 0.001$). Title dominates ($\chi^2=283.31$), followed by Sex ($=260.71$) and Pclass ($=102.88$).


8. Correlation matrix

In [112...

```

df["Sex_num"] = df["Sex"].map({"male": 0, "female": 1})
corr_cols = ["Survived", "Sex_num", "Pclass", "Age", "SibSp", "Parch", "Fare", "Fam
corr = df[corr_cols].corr()
plt.figure(figsize=(12, 9))
sns.heatmap(corr, annot=True, fmt=".2f", cmap="coolwarm", center=0, square=True,
            cbar_kws={"label": "Pearson r"})
plt.title("Feature Correlation Matrix (incl. Sex)")
plt.tight_layout()
plt.show()

```



Interpretation: This heatmap presents the Pearson correlation coefficients (r) between the main variables in the dataset. The primary objective here is to identify which factors are linearly related to survival (`Survived`).

Looking at the first row (or column) of `Survived` , we observe that **no variable has an extremely strong linear correlation (close to 1 or -1)**, confirming that **survival is not dictated by a single factor, but by a combination of several**. The key points are:

- `Sex_num` (+0.54): This is the strongest positive correlation. Being female is strongly positively correlated with survival.
- `Pclass` (-0.34): This is the strongest (negative) correlation with survival. The negative sign indicates that the higher the class number (`Pclass` 3), the lower the survival rate. This is the main marker of socioeconomic status.
- `HasCabin` (+0.32): Strong positive correlation. Simply having a registered cabin number (which is a proxy for 1st/2nd class and proximity to the lifeboat deck) significantly increases the chances of survival.
- `Fare` (+0.26): Moderate positive correlation. The more expensive the ticket, the higher the survival rate. Although related to class, it also captures nuances within the same class (luxury suite vs. standard cabin).

- **IsAlone** (-0.20): Negative correlation. Passengers traveling alone (IsAlone = 1) were at a clear disadvantage compared to those traveling with family.
- **Age** (-0.07): Very weak negative correlation. Linearly, age doesn't explain much. Because the relationship between age and survival is not linear (children were saved first, but many elderly people perished). Pearson's correlation coefficient does not capture this U-shaped relationship.
- **Family Size** (+0.02): Near-zero correlation. Here again, the relationship is not linear (families of 2 to 4 people survived better, but families of 5 or more perished). The linear correlation masks this "Goldilocks" effect (too small or too large = danger).

9. Model prediction

9.1. Advanced feature engineering

```
In [113... # Title
full['Title'] = full['Name'].str.extract(' ([A-Za-z]+)\.', expand=False)
full['Title'] = full['Title'].replace(['Lady', 'Countess', 'Capt', 'Col', 'Don', 'Dr
full['Title'] = full['Title'].replace('Mlle', 'Miss')
full['Title'] = full['Title'].replace('Ms', 'Miss')
full['Title'] = full['Title'].replace('Mme', 'Mrs')

# FamilySize
full['FamilySize'] = full['SibSp'] + full['Parch'] + 1
full['IsAlone'] = (full['FamilySize'] == 1).astype(int)

# Deck
full['Deck'] = full['Cabin'].fillna('U').astype(str).str[0]
full['HasCabin'] = full['Cabin'].notna().astype(int)

# Ticket
full['Ticket_Freq'] = full.groupby('Ticket')['Ticket'].transform('count')

# Fare
full['FarePerPerson'] = full['Fare'] / full['Ticket_Freq']

# TicketPrefix
full['TicketPrefix'] = full['Ticket'].apply(lambda x: re.sub(r'^A-Za-z', '', x).s

print("Features created")
```

Features created

9.2. Intelligent imputation and cleanup

```
In [114... # Imputation of Age by Grouped Median (Pclass & Title)
full['Age'] = full.groupby(['Title', 'Pclass'])['Age'].transform(lambda x: x.fillna
full['Age'] = full['Age'].fillna(full['Age'].median()) # Fallback

# Imputation Fare
full['Fare'] = full.groupby(['Pclass', 'Embarked'])['Fare'].transform(lambda x: x.f
```

```

full['FarePerPerson'] = full['FarePerPerson'].fillna(0)

# Imputation Embarked
full['Embarked'] = full['Embarked'].fillna(full['Embarked'].mode()[0])

# Logarithmic Transformation on Fare
full['Fare'] = np.log1p(full['Fare'])
full['FarePerPerson'] = np.log1p(full['FarePerPerson'])

# Creating age bands (Bins)
full['AgeBin'] = pd.cut(full['Age'], bins=[0, 12, 18, 35, 60, 100], labels=['Child', 'Teen', 'Adult', 'Elderly'])

# Drop unnecessary columns
full = full.drop(['Name', 'Ticket', 'Cabin', 'SibSp', 'Parch', 'PassengerId'], axis=1)

print("Data imputed and cleaned")

```

Data imputed and cleaned

9.3. Encoding of categorical variables

In [115...

```

# Encoding One-Hot
categorical = ['Sex', 'Embarked', 'Title', 'Deck', 'TicketPrefix', 'AgeBin']
full = pd.get_dummies(full, columns=categorical, drop_first=True)

# Separation Train / Test
train_proc = full[full['__set'] == 'train'].drop('__set', axis=1)
test_proc = full[full['__set'] == 'test'].drop(['__set', 'Survived'], axis=1)

# Column alignment
test_proc = test_proc.reindex(columns=train_proc.columns, fill_value=0)
test_proc = test_proc.drop('Survived', axis=1)

X = train_proc.drop('Survived', axis=1)
y = train_proc['Survived'].astype(int)
X_test_final = test_proc

print(f"Shape X_train: {X.shape} | Shape X_test: {X_test_final.shape}")

```

Shape X_train: (891, 59) | Shape X_test: (418, 59)

9.4. Modeling and training

In [116...

```

# Random Forest
test_ids = test["PassengerId"]

rf = RandomForestClassifier(
    n_estimators=500, max_depth=6, min_samples_leaf=2,
    max_features='sqrt', random_state=42, n_jobs=-1
)

# Gradient Boosting
gb = GradientBoostingClassifier(
    n_estimators=200, learning_rate=0.05, max_depth=4,

```

```

        min_samples_leaf=2, random_state=42
    )

    # Vector Machine Support
    svm_pipe = Pipeline([
        ('scaler', StandardScaler()),
        ('svm', SVC(probability=True, C=1.0, kernel='rbf', random_state=42))
    ])

    # Logistic regression
    logreg_pipe = Pipeline([
        ('scaler', StandardScaler()),
        ('logreg', LogisticRegression(C=0.1, solver='liblinear', random_state=42))
    ])

    # Creation of the assembly model
    ensemble_model = VotingClassifier(
        estimators=[
            ('rf', rf),
            ('gb', gb),
            ('svm', svm_pipe),
            ('logreg', logreg_pipe)
        ],
        voting='soft',
        weights=[2, 3, 1, 1]
    )

    print("Training of the ensemble model on the entire dataset...")
    ensemble_model.fit(X, y)
    print("✅ Training complete")

    predictions = ensemble_model.predict(X_test_final)

    submission = pd.DataFrame({
        'PassengerId': test_ids,
        'Survived': predictions
    })

    submission.to_csv('submission.csv', index=False)

    print("\n📁 'submission.csv' file generated successfully")
    print("Predictions overview :")
    print(submission.head())

```

Training of the ensemble model on the entire dataset...

✅ Training complete

📁 'submission.csv' file generated successfully

Predictions overview :

	PassengerId	Survived
0	892	0
1	893	0
2	894	0
3	895	0
4	896	1

Conclusion

Through this comprehensive analytical journey, from initial univariate explorations to complex predictive modeling, we have uncovered the critical determinants of survival aboard the Titanic.

The data clearly demonstrates that survival was not merely a game of chance. Distinct passenger profiles, heavily influenced by socio-economic class, sex, and age, were significantly more likely to be spared from the tragedy.

By quantifying these determining factors, our machine learning models successfully translated historical observations into strong predictive patterns. This EDA not only highlights the stark realities of the disaster but also showcases the power of data analysis and data science in extracting meaningful narratives from raw, historical records.